

Intro to Databases

SQL

What is a database?

A database **persists** information
and is **accessible** via code

```
graph TD; A[accessible] --- B[organized]; A --- C[queryable]; A --- D[manageable]
```

organized **queryable** **manageable**

Table / Relation

Column / Attribute / Field Column / Attribute / Field Column / Attribute / Field

	ID	Name	Type
Row / Tuple / Instance	1	Pikachu	lightning
Row / Tuple / Instance	2	Squirtle	water
Row / Tuple / Instance	3	Charmander	fire
Row / Tuple / Instance	4	Bulbasaur	grass



SQL

```
-- Pikachu, I choose you!  
SELECT id, name  
  FROM pokemon  
 WHERE type = 'lightning'  
LIMIT 1
```

ACID Guarantees

- **Atomicity**
- **Consistency**
- **Isolation**
- **Durability**

Atomicity - A set of database operations that must occur together. All or nothing. A transaction cannot partially finish, it must either fail or complete.

Consistency - If a process has a writer, no other process can read from it, and no other process can write to it, which gives us consistent information at all times.

Isolation - Multiple clients can make queries to read and update without the risk of deadlock or starvation.

Durability - Store information without power (flash drive / hard drive, for example doesn't need power and still stores data)

Atomic Transactions

- **atomic transaction: A set of database operations that must occur together**
 - i.e. A debit to one bank account, and a credit to another
- **A transaction must either succeed or fail; it cannot partially complete.**
- **Every database query is represented by a transaction**

We might also say that a database is atomic or a database has the property of atomicity.

We can define where transactions start and end. You may need to as a programmer

Consistency

- **Specify rules that columns need to follow**
 - Gender column can only contain M, F, or U.
 - Savings account must start with S or checking with C
 - Column cannot be null
- **Protect the database from inconsistencies and simplify software logic**
 - Allows software to make assumptions about underlying data

Resource Management

- Processes can be readers and writers
- Files can have many readers
- If a process has a writer, no other process can read from it, and no other process can write to it

Proposed File Scheme

- Suppose that we have decided not to use a database and instead store our data in a series of files.
- How might our setup fail to serve queries from multiple users?

Deadlock



- Three files: A, B, and C
- Process 1 needs files A then B for writing
- Process 2 needs files B then C for writing
- Process 3 needs files C then A for writing
- What happens if all three processes start a database request at once?

Databases give us concurrency (Isolation)

- Multiple clients can make queries to read and update without the risk of deadlock or starvation.

Persistence/Durability

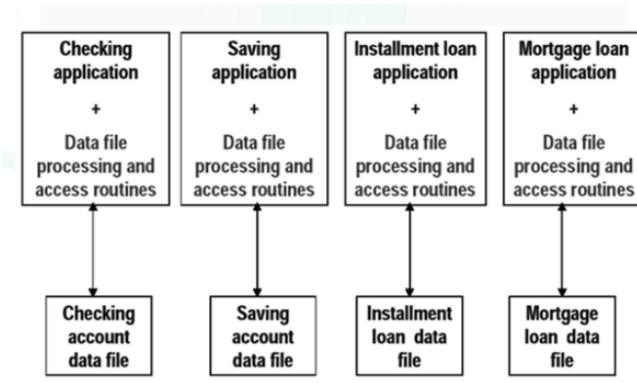
- Files are also persistence (store information without power)

- What happens when two people try to access the same file at the same time?
- the second process waits and retries until the first process finishes?

How Did We End Up Here?

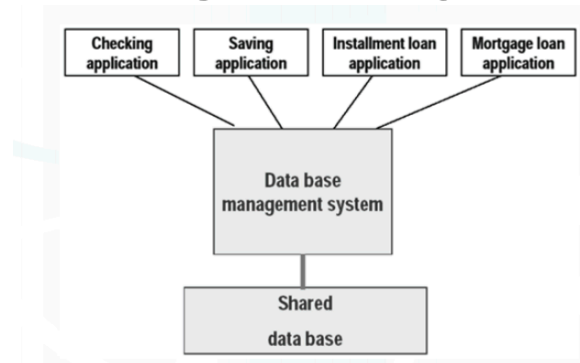
Before Relational DBs (ca. < 1970s)

- Data stored in custom “data files”
- Queried via application-specific code
- **Advantages**
 - Middle layer not needed
 - Solutions customized for each application
- **Disadvantages**
 - Hard to change the system
 - Knowledge not compounding
 - Data-transfer is difficult



- Separate file for every topic. Routines and functions for dealing with each particular file.
- But every company formatted their files differently based on developer whims and needed their own accessor functions.
- Making changes was really hard seeing as formats were inflexible and etched in stone.

Database Management Systems (DBMS)




- One layer and language to store and access data
- Sold as a way for “non-technical people” to manage data

Ubiquitous
Standardized
Don't need to be a programmer
common way of thinking about data. And a common language for querying it.

“Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation).”

– E. F. CODD,
A RELATIONAL MODEL OF DATA FOR
LARGE SHARED DATA BANKS

Relational Databases & Logic



- 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- Wrote Alpha (never implemented) as a *query language*
- IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created **Structured English Query Lang**
- **SEQUEL** way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- **SQL became the standard (ANSI 1986, ISO 1987)**
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*

Appreciating Databases

- Ubiquitous
- Standardized
- Complex / deep
- Powerful: database admins are
 - Feared by developers
 - ...but also taken for granted until things break
 - Befriended by business people
 - Contacted by the government for secret data (e.g. NSA)

Progression of Databases

- **Navigational (< 1970s)**
 - More common during tape era; entries had references to next entries.
- **Relational (> 1970s)**
 - Based on relational (table-based) logic, see E.F. Codd.
- **NoSQL (> 2000s)**
 - "Not only SQL" — document storage, for example.

Navigational: common during tape era, entries had references to the next entries.

Relational: based on relational (table-based) logic, EF Codd.

NoSQL ("not only SQL"): document storage.


RDBMS vs NoSQL

- **A DBMS doesn't *have* to be relational**
 - Remember, DBMS is just an application that intelligently stores data and can answer requests to manage that data
- **Lately, many "NoSQL" or non-relational DBMSs have been gaining popularity**
 - Graph databases (e.g. Neo4j)
 - Document databases (e.g. MongoDB)
 - Hybrids (e.g. PostgreSQL)
- **RDBMSs still remain the #1 DB option for now**



Some well-known rDBMSs





PostgreSQL

Why PostgreSQL?

- Advanced, powerful, and popular
- Rapid open source development
- Highly extensible (stored procedures)
- Deep SQL standards compliance
- NoSQL ("Not Only SQL"), objective support
- Excellent transactions / ACID reliability; focus on integrity
- Multi-user management / administration

History of PostgreSQL

- 1970s at UC Berkeley:
Interactive **G**raphics **R**etrieval **S**ystem (INGRES)
- 1980s: POSTGRES ("Post-Ingres")
- 1995: POSTQUEL and Postgres95.
 - monitor -> psql
- 1996: Adopted by the open source community
 - Ongoing: stability, testing, documentation, new features
 - PostgreSQL

psql

```
psql [local]:5432 glebec # - \l
List of databases

```

Name	Owner	Encoding	Collate	Ctype	Access privileges
assessmentexpresssequelize	fullstack	UTF8	en_US.UTF-8	en_US.UTF-8	
author	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
checkpoint-angular	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
checkpoint-express-review	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
glebec	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
joke	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
postgres	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
sequelizecheckpoint	fullstack	UTF8	en_US.UTF-8	en_US.UTF-8	
template0	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	sc/glebec
template1	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	sc/glebec
tripplanner	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	glebec=CTC/glebec
twitterdb	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	
wikistack	glebec	UTF8	en_US.UTF-8	en_US.UTF-8	

```
(13 rows)

psql [local]:5432 glebec # - \c author
You are now connected to database "author" as user "glebec".

psql [local]:5432 glebec # author \dt
List of relations

```

Schema	Name	Type	Owner
public	stories	table	glebec
public	users	table	glebec

```
(2 rows)

psql [local]:5432 glebec # author
```



pgcli

```
stayupdated_test> \d
+-----+-----+-----+-----+
| Schema | Name          | Type          | Owner  |
+-----+-----+-----+-----+
| public  | admins        | table         | amjith |
| public  | cpes          | table         | amjith |
| public  | goose_db_version | table        | amjith |
| public  | goose_db_version_id_seq | sequence    | amjith |
| public  | packages      | table         | amjith |
| public  | packages_id_seq | sequence     | amjith |
| public  | users         | table         | amjith |
| public  | users_id_seq  | sequence     | amjith |
| public  | vulnerabilities | table        | amjith |
| public  | vulnerabilities_cpes | table      | amjith |
| public  | vulnerabilities_id_seq | sequence    | amjith |
+-----+-----+-----+-----+

SELECT 11
stayupdated_test> SELECT * FROM users;
+-----+-----+-----+-----+-----+
| id | display_name | password | email | created_on |
+-----+-----+-----+-----+-----+
| 177 | DisplayName1 | 1024cms | user@ex.com | 2014-11-15 15:02:50.094560 |
| 180 | testname2    | pas5w0rd | email@ex.com | 2014-11-28 10:25:46.170660 |
| 181 | amjith       | password | amjith@amjith.amjith | 2014-11-28 18:39:48.195067 |
+-----+-----+-----+-----+-----+

SELECT 3
stayupdated_test> SELECT * FROM
admins
cpes
goose_db_version
packages
users
```

Postico

ReportingreportingforexConnectedPostgreSQL 9.4.5

SQL Query

currencies

daily_proceeds

forex

forex_sources

fs_daily_proceeds

fs_daily_sales

fs_orders

fs_proceeds_weekly

fs_weekly_sales

itc_daily_proceeds

itc_daily_sales

itc_proceeds_weekly

itc_reports

itc_reports_daily

itc_reports_monthly

itc_reports_weekly

Search

currency	base	rate	date	source_id
AED	USD	3.67291	2014-07-25	1
AFN	USD	56.485726	2014-07-25	1
ALL	USD	103.5838	2014-07-25	1
AMD	USD	410.086	2014-07-25	1
ANG	USD	1.787	2014-07-25	1
AOA	USD	96.952626	2014-07-25	1
ARS	USD	8.169642	2014-07-25	1
AUD	USD	1.062844	2014-07-25	1
AWG	USD	1.79	2014-07-25	1
AZN	USD	0.784067	2014-07-25	1
BAM	USD	1.453024	2014-07-25	1
BBD	USD	2	2014-07-25	1
BDT	USD	77.61971	2014-07-25	1
BGN	USD	1.452543	2014-07-25	1

Content

Structure

+

Filter

Page 1 of 342

currency

MULTIPLE

base

USD

rate

MULTIPLE

date

2014-07-25

source_id

forex_source:

id

name

1

Open

Exchange R...

LAB